

Apache Spark 2.0 Catalyst Engine

Preliminary Notes

By Oliver Tupran @ Lunatech 2016

Intro

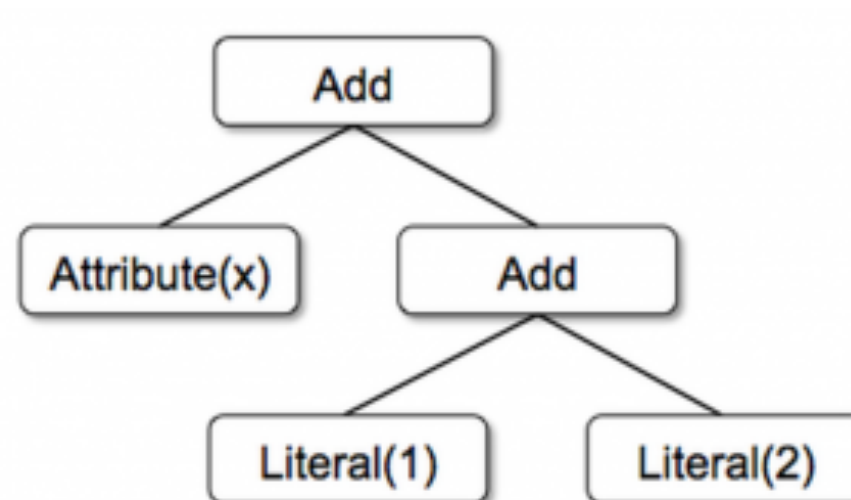
- Catalyst is a query optimisation engine
- The Catalyst engine was introduced in 2014 in Spark 1.0.0 (see [SPARK-1251](#))

Trees

- `TreeNode` is the main data type in Catalyst
- `TreeNode`s are immutable
- `TreeNode`s can be manipulated through recursive functional transformations (e.g. `map`, `flatMap`, `foreach`, `transform...`)
- `TreeNode` has two main implementations:
`Expression` and `QueryPlan`

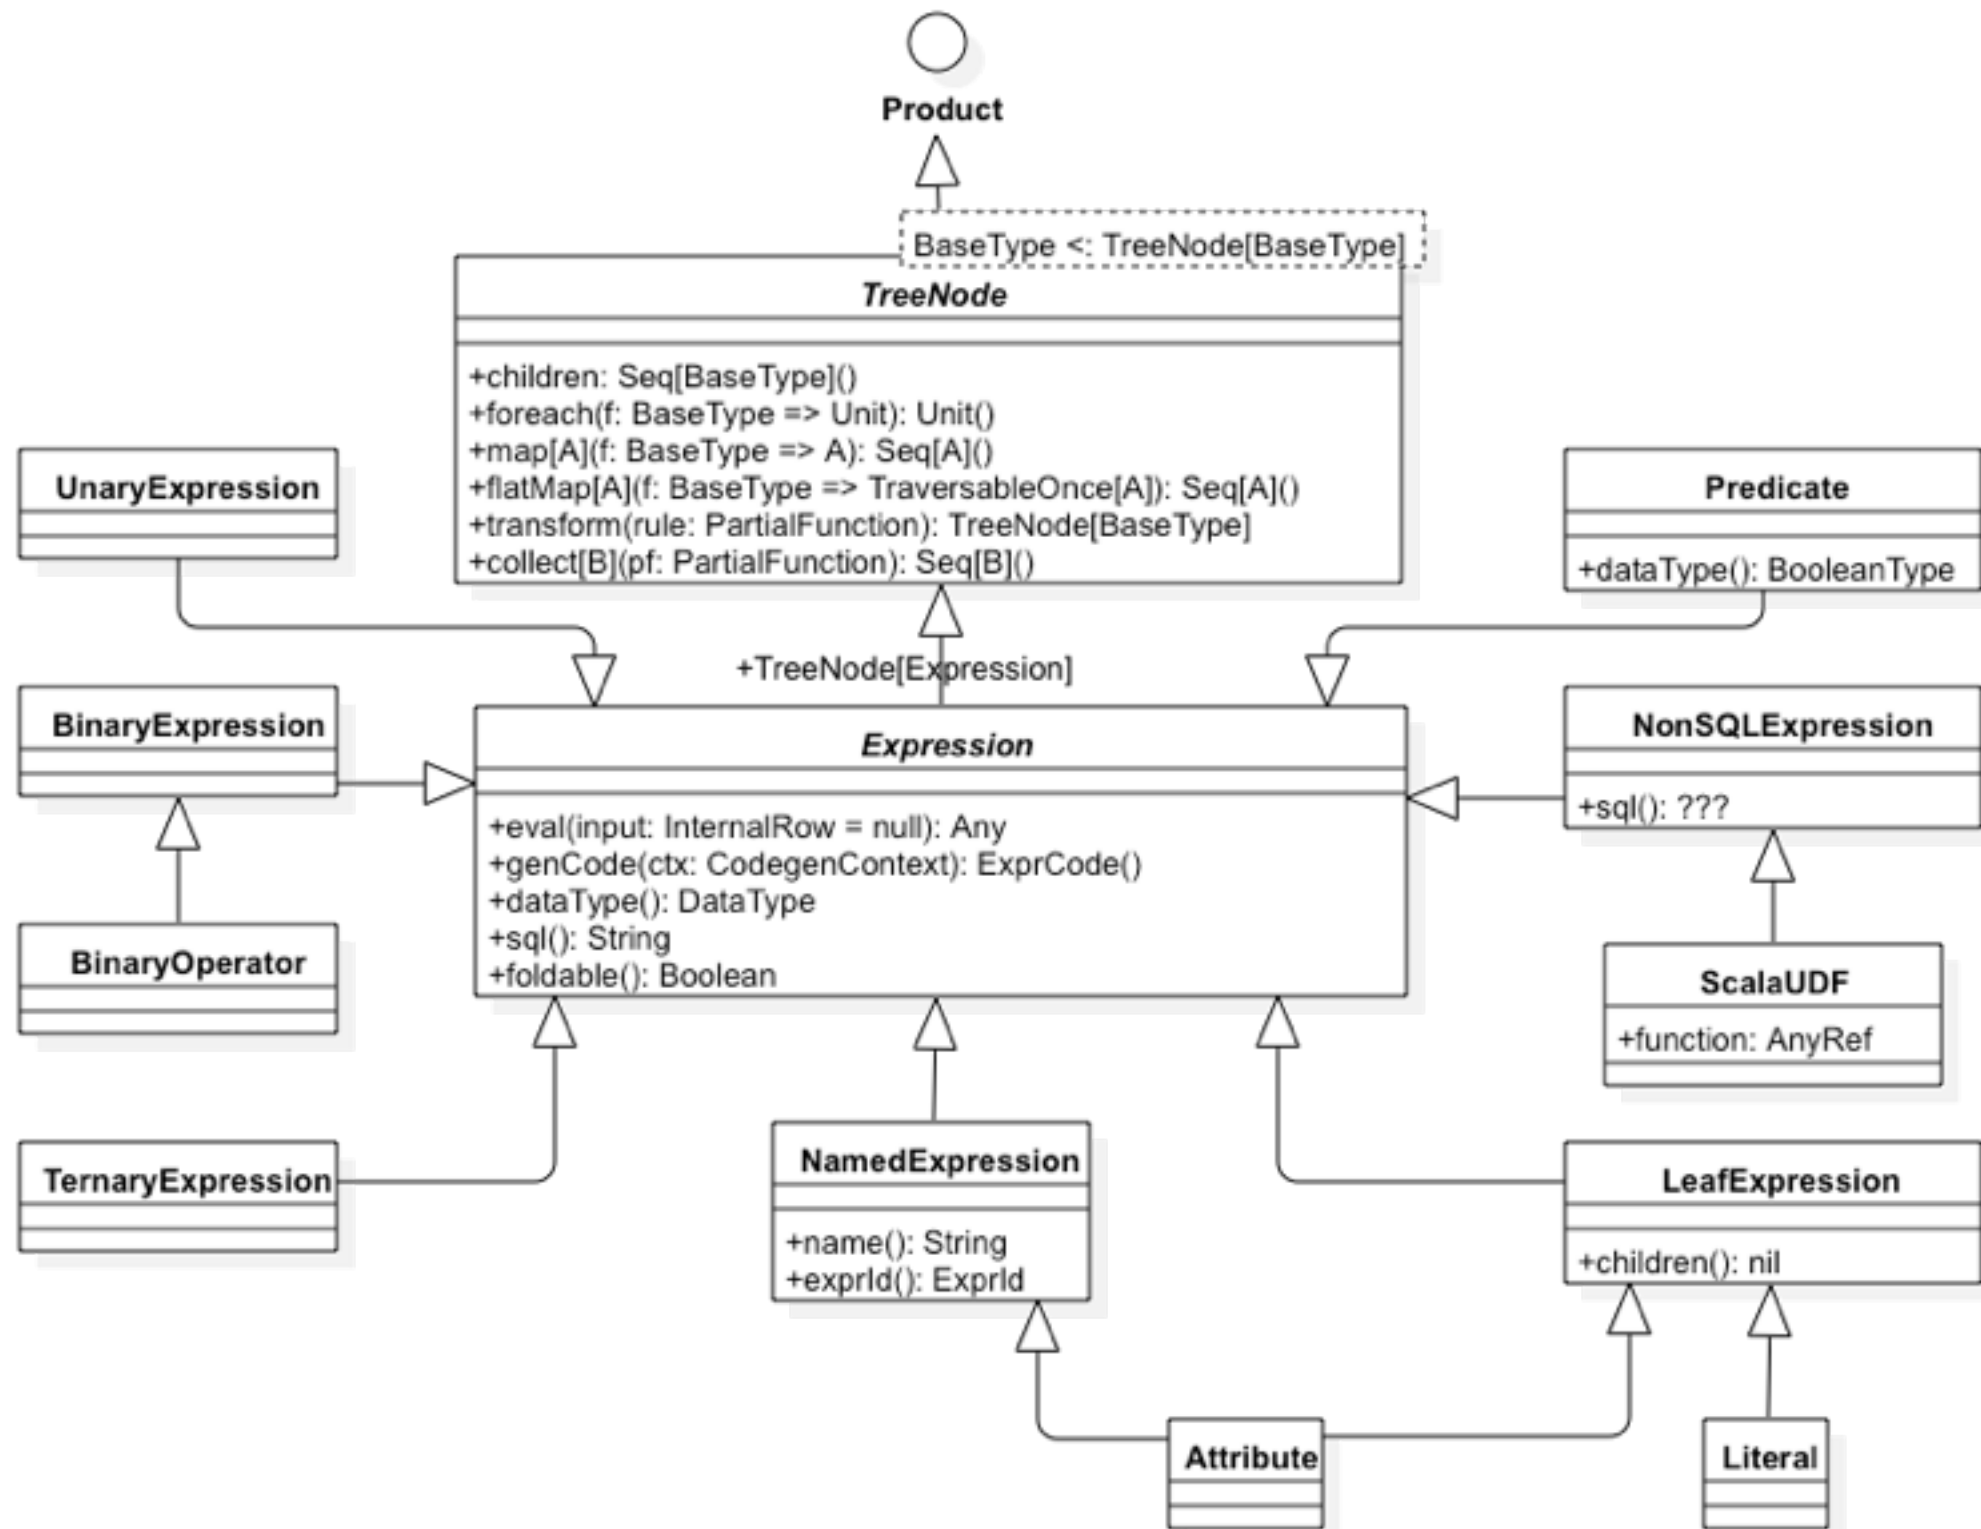
Expression Example

`x + (1 + 2)`



`Add(Attribute("x"), Add(Literal(1), Literal(2)))`

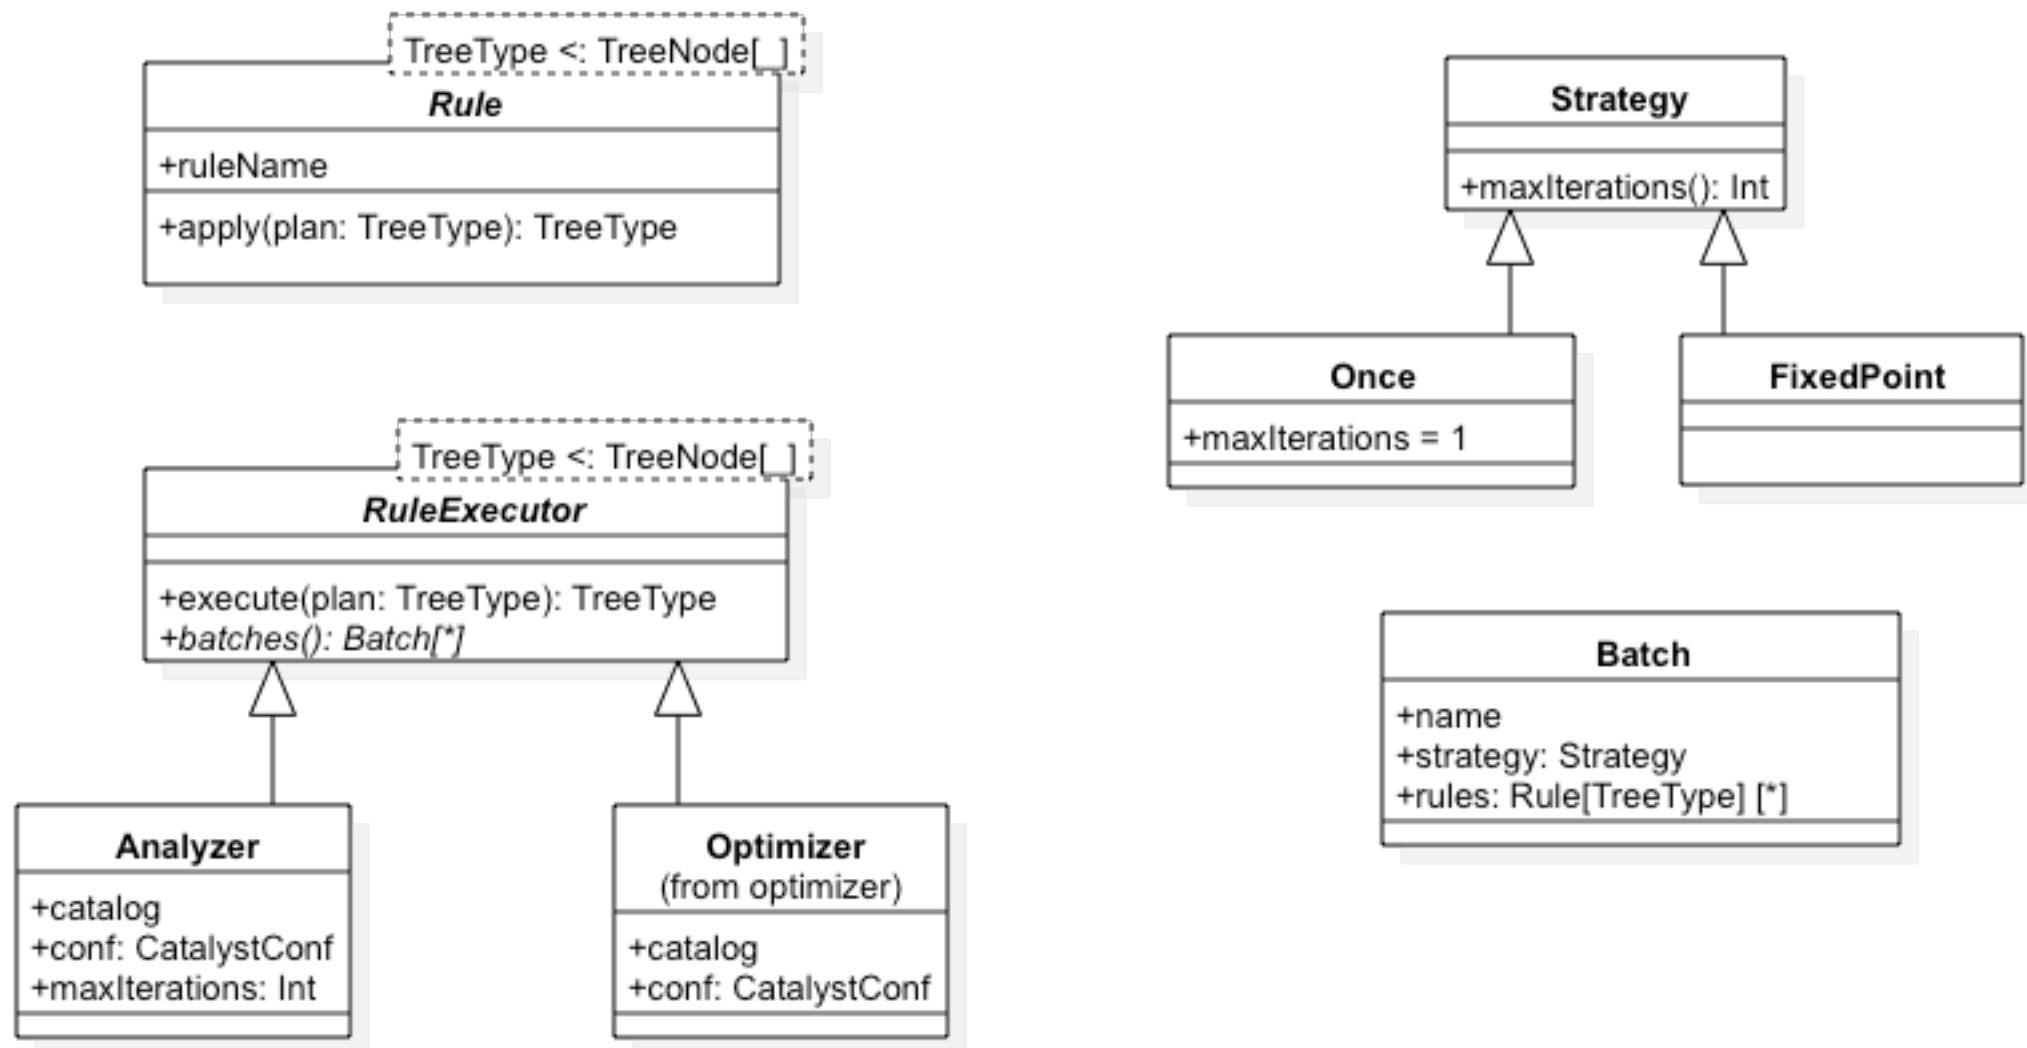
Expressions



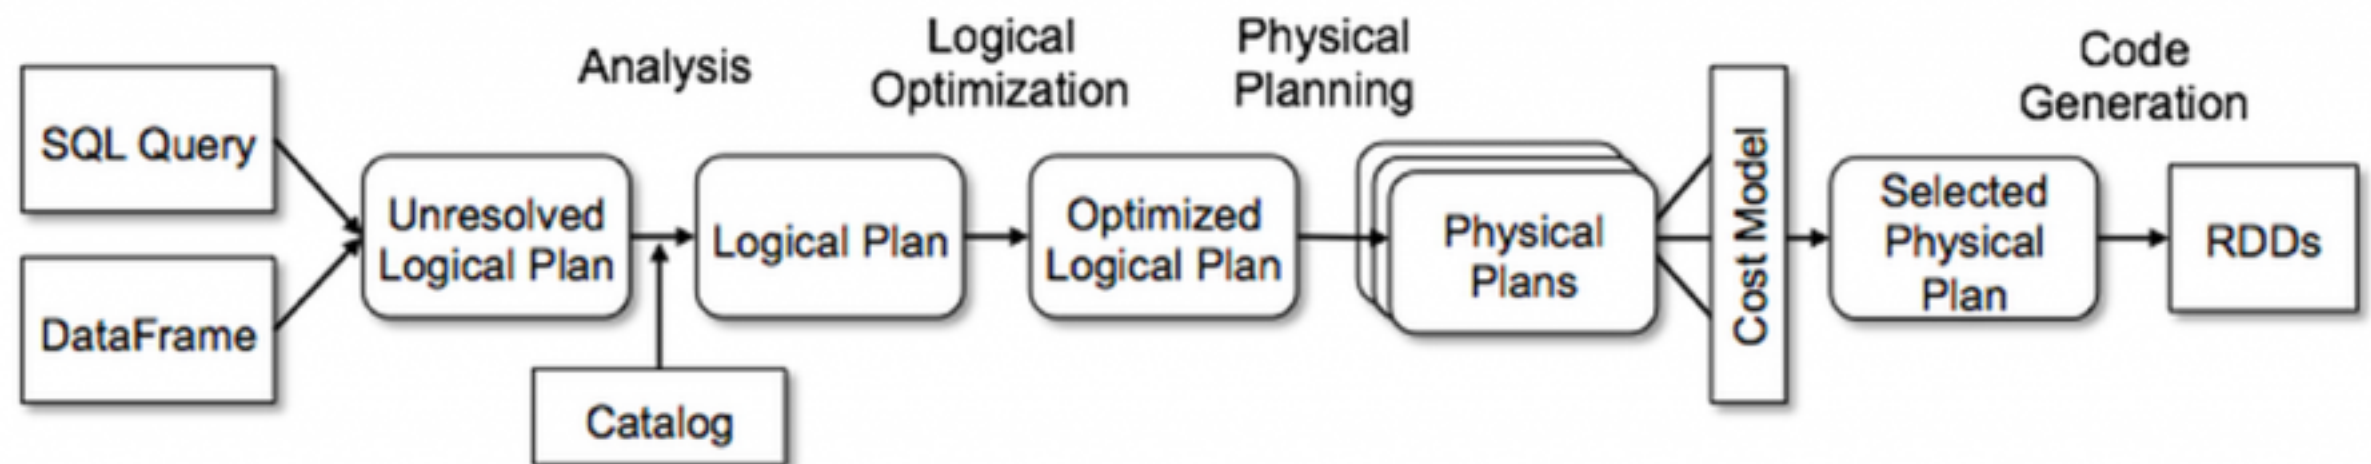
Rules

- Rules are functions, or partial functions, that transform a `TreeNode` into a `TreeNode` (though `Rule` does not extend `Function1`)
- Rules are applied using `RuleExecutors`
- `RuleExecutors` apply batches of Rules to the `TreeNode` using various strategies (most common are `Once` or `FixedPoint`)
- Two main implementations of `RuleExecutor` are `Analyser` and `Optimizer`

Rules



Catalyst Work Flow

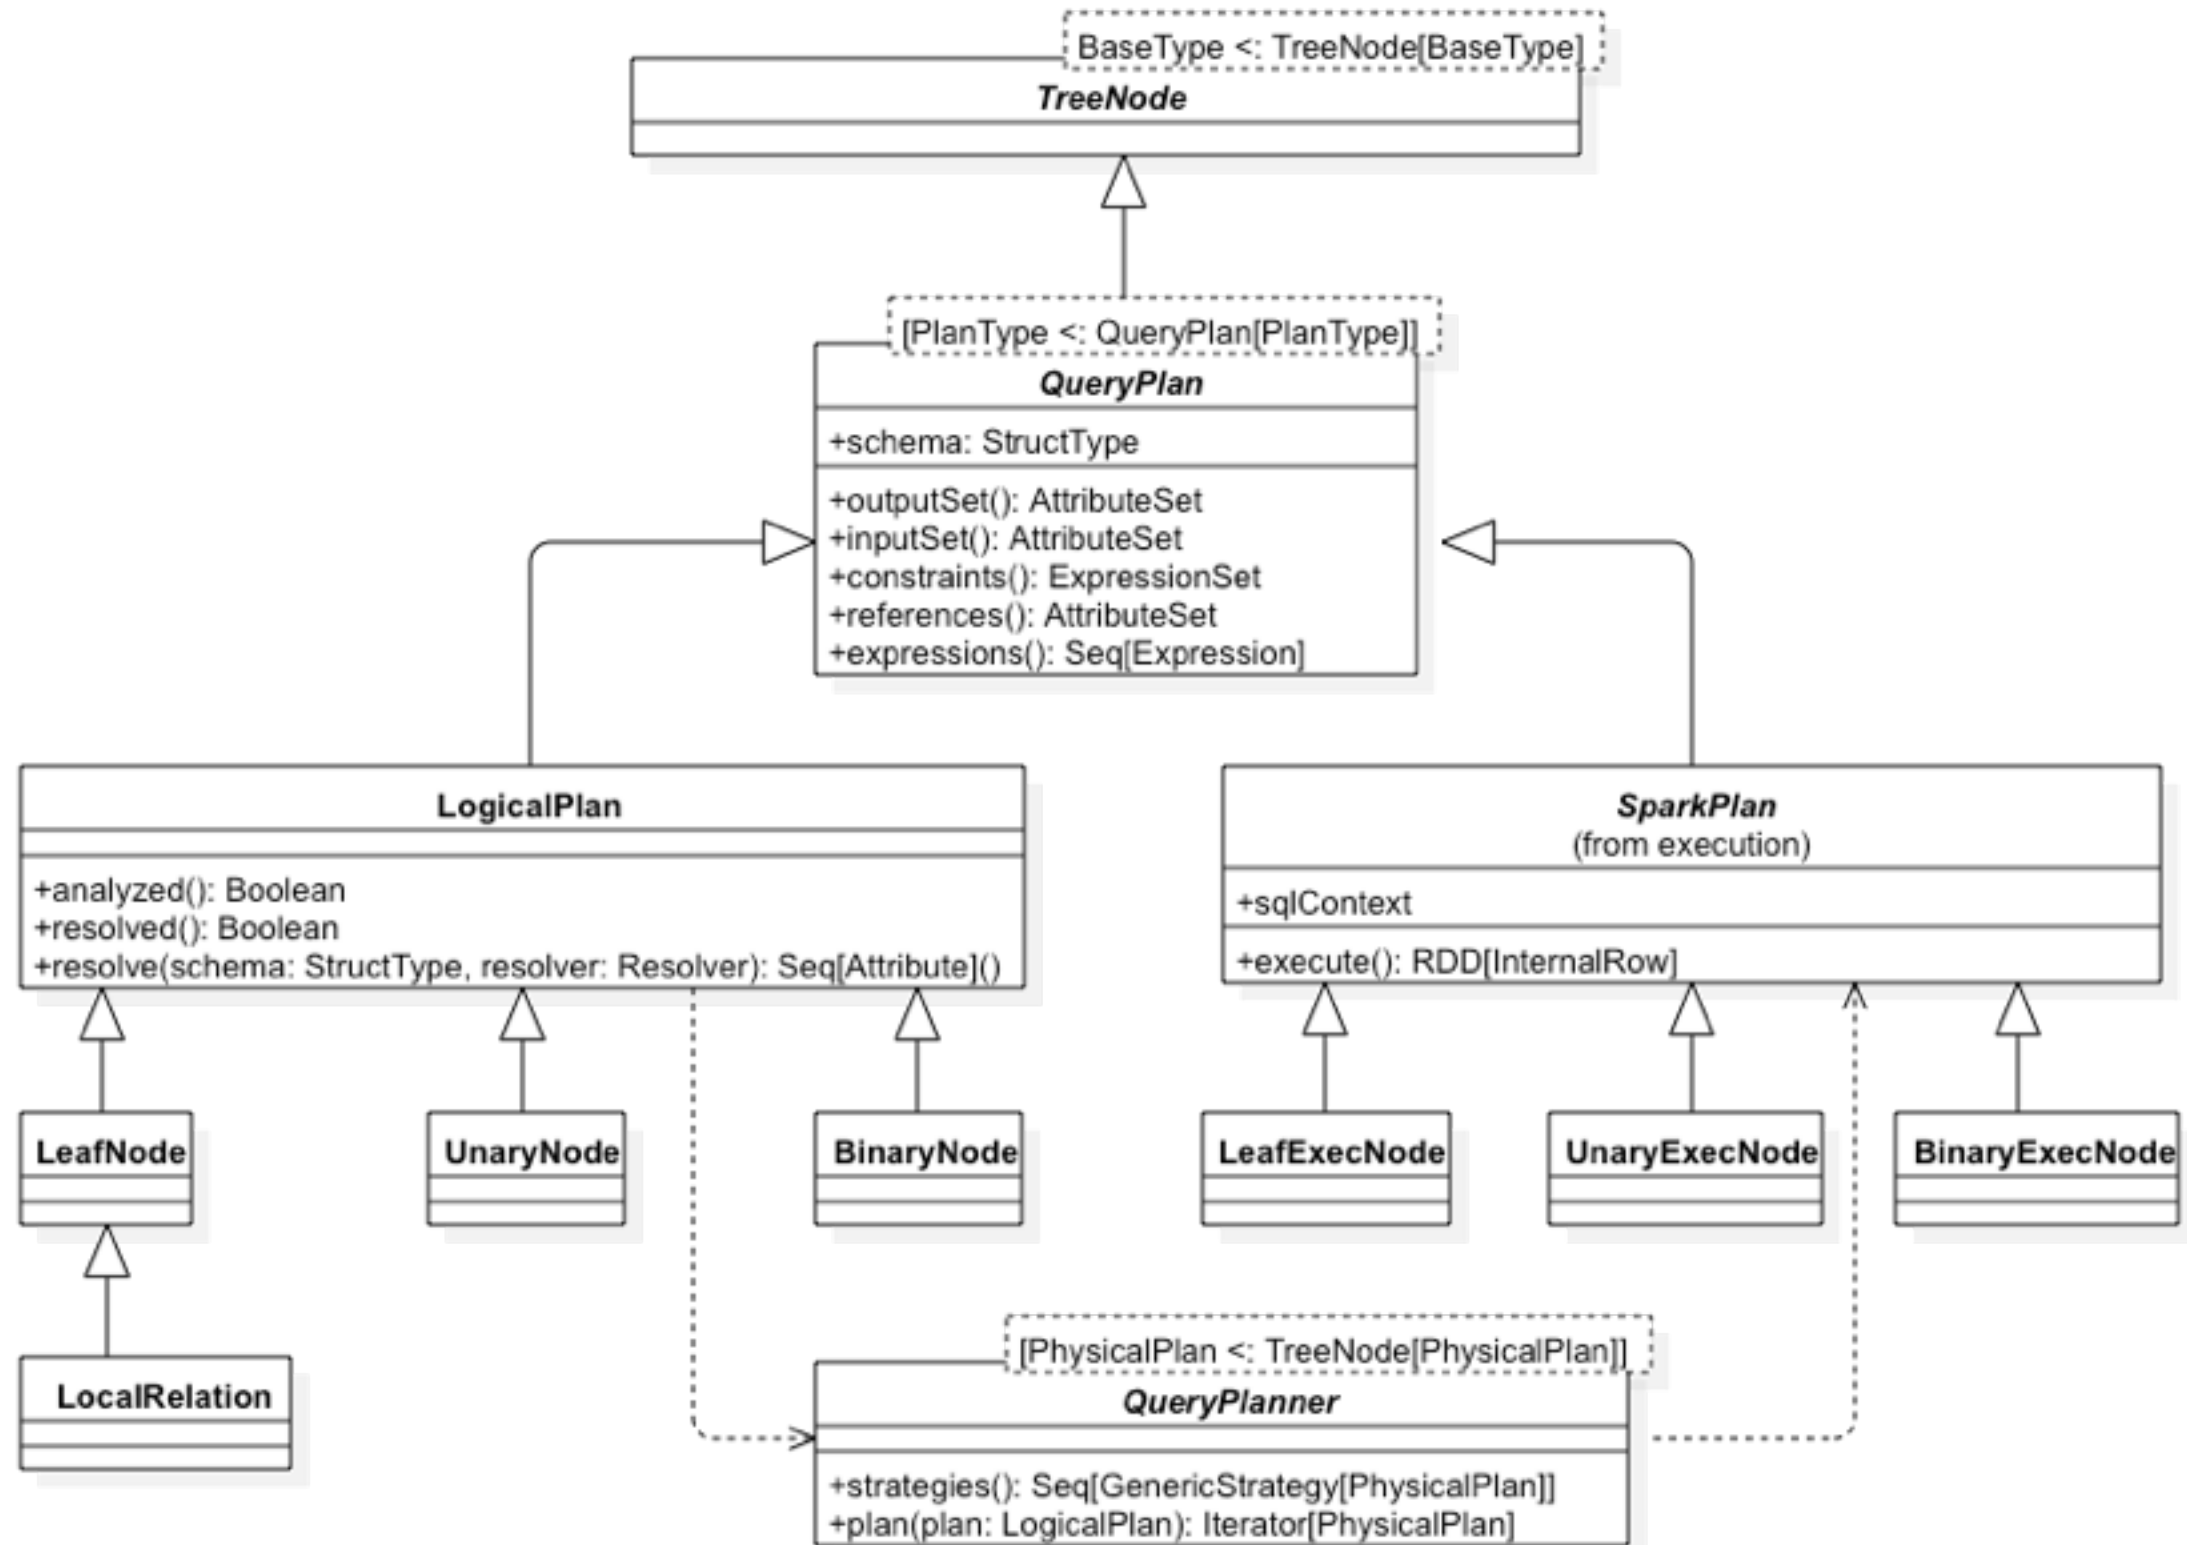


Analysis

- Transform raw, unresolved `LogicalPlan` into a resolved and analysed `LogicalPlan`
- Steps
 - resolve relations (tables)
 - resolve attributes
 - resolving attributes referring to the same value
 - resolve attribute types and casts

See [Analyzer](#)

Query Plans



Logical Optimisation

- The logical optimiser is producing a new logical plan
- Examples:
 - constant folding
 - predicate pushdown
 - projection pruning
 - null propagation
 - boolean expression simplification

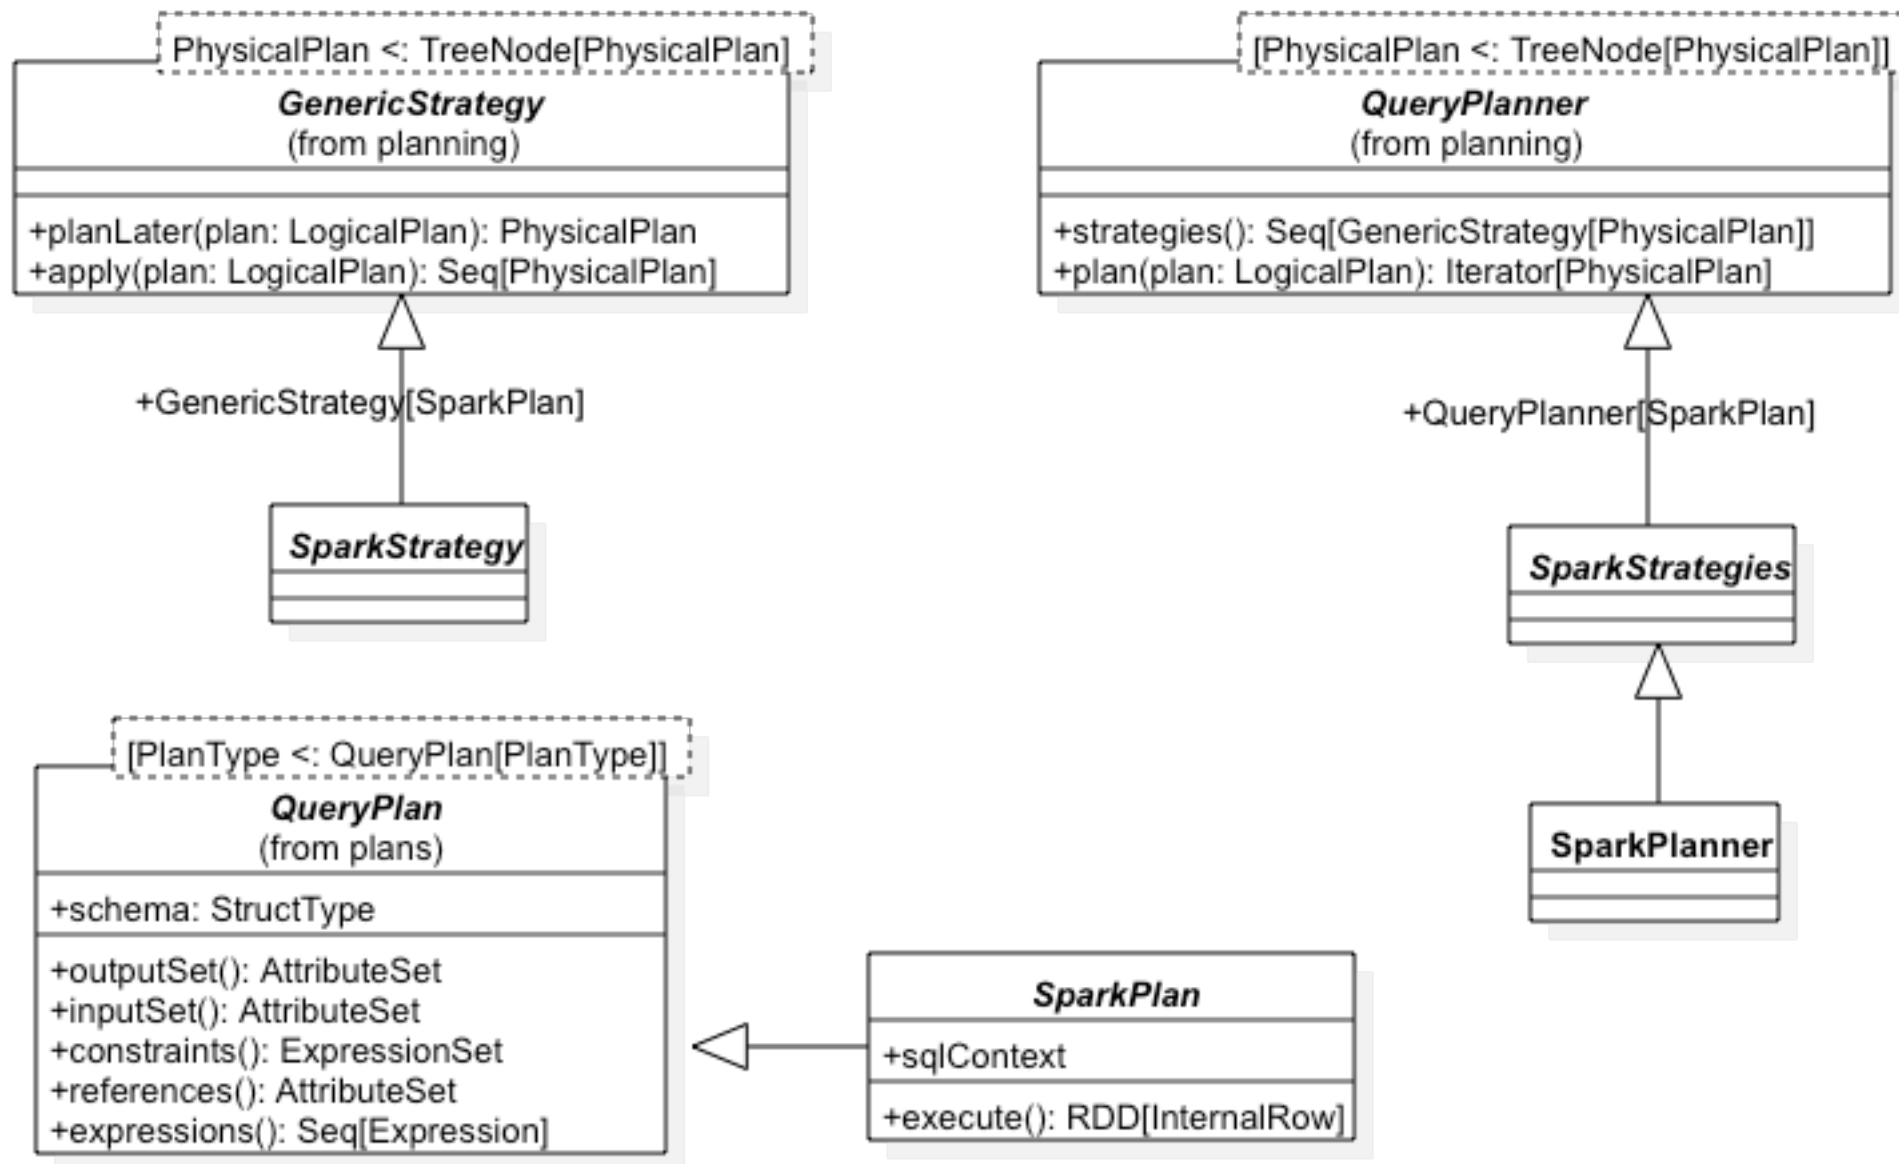
See [Optimizer](#)

Physical Planning

- From a `LogicalPlan` produce one or more `PhysicalPlans`
- On the resulted `PhysicalPlans` both cost-based and rules-based optimisations are applied
- If possible it will do a Projection Push-Down, to the external system (e.g. HDFS, Cassandra, RDMS...)
- The resulted `PhysicalPlan` is optimised for Spark

See [SparkStrategies](#)

Execution



Code Generation

- Mainly based on Scala Quasiquotes, but there is some string based generation as well
- No virtual function calls, linear code
- Entire code fits in one function => data passes through CPU registers vs memory

See CodeGenerator

TODOs

- Find if Catalyst truly needs a dependency to spark-core
- Find the best way to deliver the code:
 - source or
 - byte code or
 - logical plan
- Test more complex plans

References

- Spark SQL: Relational Data Processing in Spark
- SPARK-1251 Catalyst
- SPARK-12795 Whole stage codegen
- Deep Dive into Spark SQL's Catalyst Optimizer
- Apache Spark as a Compiler
- anatomy-of-spark-catalyst